

ECE 421
Assignment 1 Design Report

By: Vishal Patel, Jose Ramirez, Rizwan Qureshi, Curtis
Goud, Jori Romans

Design Rationale:

We designed our sparse matrix library with the goal of providing computational optimization for sparse matrices when performing matrix operations. We chose to support:

- Basic scalar and vector matrix operations
- Custom iterators that allow you to parse specific columns or rows.
- Support for “list of lists”, “dictionary of keys” and the new “Yale” storage types
- Conversion between all matrix storage types
- LU decomposition and Cholesky decomposition
- Construction and checking of special matrix types: Zero matrix, Identity matrix, Tridiagonal matrix

This library has a “list of lists” and “dictionary of keys” implementation for storage of matrices where efficient modification is required and a new “Yale” implementation for storage of matrices where sufficient access and matrix operations is required. We plan to create an abstract factory to create matrix factories for each storage type. Our SMatrix library will contain the contracts and a general implementation of matrix operations and construction using the delegate design pattern which will delegate matrix operations to the respective implementation for that storage type.

To supplement our design, we chose to use the NMatrix gem that supports a large amount of features which include but not limited to: shorthand notation for matrix creation, list of lists and yale matrix storage, multiple data types, many element wise and matrix operations and custom iterators for Yale matrices. The main advantage to NMatrix is that it supports most of the features we chose to implement which will save us development time. The disadvantage to NMatrix is that a couple important matrix solving functions such as Eigenvalue decomposition is not available (therefore we may have to implement it ourselves).

System Research

1. What is a sparse matrix and what features should it possess?
 - A matrix is sparse if most of the elements are zero.
 - Sparsity of a matrix = $(\# \text{ of 0-valued elements}) / (\text{total } \# \text{ of elements}) = 1 - (\text{density of a matrix})$. A matrix is “sparse” if sparsity > 0.5
2. What sources of information should you consult to derive features? Do analogies exist? If so with what?
 - Analogy 1: Consider a line of balls connected by strings. If each adjacent ball was connected by a spring then the system is sparse. If each ball had a spring connecting it with all other balls then the system is dense.
 - Sources of Information:
 - Consulting Wikipedia (https://en.wikipedia.org/wiki/Sparse_matrix)
 - Textbook: Selected Chapters from Elementary Linear Algebra, Eleventh Edition by: Anton & Rorres
 - Existing linear algebra library: <https://github.com/SciRuby/nmatrix> . Has an experimental implementation in Ruby.
 - Existing implementations of Sparse Matrix libraries - <https://www.mathworks.com/help/matlab/math/sparse-matrix-operations.html#f6-14457>
 - Default matrix library for Ruby: <https://docs.ruby-lang.org/en/2.4.0/Matrix.html>
3. Who is likely to be the user of a sparse matrix package? What features are they likely to demand?
 - Sparse matrices correspond to loosely coupled systems and are used in machine learning, combinatorics and network theory which have a low density of significant data. Additionally, large sparse matrices appear in scientific and engineering applications when solving partial differential equations.
 - For machine learning, sparse matrices are usually used to store large quantities of training data. A user for this purpose would likely want efficient storage and iteration of a large sparse matrix.
 - For computer graphics, It is many operations of matrices multiplying themselves.
 - For Google or Amazon, everytime they recommend you something they multiply sparse matrices to yield that result for you.
 - For Google’s Adsense, they calculate positive eigenvectors to determine which advertisements to show you.
 - For information retrieval, the index is essentially an intelligently crafted sparse matrix.

- Other features that are common to other matrix libraries and therefore likely to be demanded are:
 - Basic matrix scalar and vector operations including: Add, subtract, multiply, transposition, trace, exponents, inverse, determinants,
 - Slice assignments, enumerators for matrices
 - Matrix and vector storage in multiple formats
 - Interconversion between storage and data types
 - Matrix slicing by copy and reference
 - Checking tridiagonal, identity, symmetric
 - LU and Cholesky factorization

4. What is a tri-diagonal matrix?

- A tri-diagonal matrix is a matrix such that all non zero elements exist on the main diagonal, the diagonal existing one below the main diagonal, and one above the main diagonal.

$$f_n = \begin{pmatrix} a_1 & b_1 & & & \\ c_1 & a_2 & b_2 & & \\ & c_2 & \ddots & \ddots & \\ & & \ddots & \ddots & b_{n-1} \\ & & & c_{n-1} & a_n \end{pmatrix}.$$

5. What is the relationship between a tri-diagonal matrix and a generic sparse matrix?

- Tri-diagonal matrices share the same property with sparse matrices which is that they are filled with many zero elements and relatively few non-zero elements.
- A very efficient structure for the diagonal matrix is to store just the entries in the main diagonal as 3 one-dimensional arrays, so a diagonal $n \times n$ matrix requires only $3n$ entries.
- NOTE: Not all tri-diagonal matrices are sparse matrices.

$$\begin{pmatrix} 1 & 7 & 0 & 0 \\ 2 & 3 & 8 & 0 \\ 0 & 5 & 6 & 4 \\ 0 & 0 & 9 & 10 \end{pmatrix}$$

Tri Diagonal Matrix

This matrix has a sparsity of $6/16 = 0.375 < 0.50$ which is the threshold to be considered a sparse matrix.

6. Are tri-diagonal matrices important? And should they impact your design? If so, how?
 - Yes, they are important when performing matrix operations. When using a tri-diagonal matrix, we can efficiently process and calculate matrix operations such as determinants, inverses, transpositions etc. compared to normal matrices.
 - This will impact our design as recognizing these special cases will allow us to apply optimizations in processing these matrices, especially for larger ones.

7. What is a good data representation for a sparse matrix?
 - We can have substantial memory reductions by storing only the non-zero elements. It is not possible to achieve both efficient modification and efficient access time at once.
 - If we want efficient modification, we could use:
 - Dictionary of keys - Maps (Row, column) key to a value. Elements missing from the dictionary are zero. Good for constructing a sparse matrix in random order but poor for iterating over.
 - List of lists - Stores one list per **row** with each entry containing **column** index and the value. Typically sorted by column index. (The bolded words refer to the original matrix.)
 - Coordinated list - Stores a list of (row, column, value) tuples. Typically sorted by row index then column index to improve random access time.
 - If we want efficient access and matrix operations:
 - Compressed sparse row (or Yale method) - Stores a sparse matrix M in 3 arrays denoted A , JA and IA . The array A , holds all nonzero entries in M in row-major order. The array JA stores the column index in M for each element in A . Note that $\text{Length}(A) == \text{Length}(JA) == \# \text{ of nonzero entries}$. Finally, the array IA , is defined recursively:
 - $IA[0] = 0$
 - $IA[i] = IA[i - 1] + (\text{number of nonzero elements on the } i\text{-th row in the original matrix})$
 - Compressed sparse column is similar to above.

8. Assume that you have a customer for your sparse matrix package. The customer states that their primary requirements as: for a $N \times N$ matrix with m non-zero entries
 - Storage should be $\sim O(km)$, where $k \ll N$ and m is any arbitrary type defined in your design.
 - Adding the $m+1$ value into the matrix should have an execution time of $\sim O(p)$ where the execution time of all method calls in standard Ruby container classes is considered to have a unit value and $p \ll m$ ideally $p = 1$

In this scenario, what is a good data representation for a sparse matrix?

- We should use both the “list of lists” and the new “Yale” representation. The “list of lists” implementation has a storage complexity of $O(2m)$ which matches the above constraint for memory space. It also has an insertion time of $O(1)$ which complies with the execution constraint of $O(p)$. The new “Yale” implementation allows us to perform matrix operations more efficiently with a storage complexity of $O((2 \times NZV) + m + 1)$ and it costs $O(1)$ to look up a row in a Yale matrix with $O(\log(n))$ to look up an entry within a row.

9. Design Patterns are common tricks which normally enshrine good practice

- Explain the design patterns: Delegate and Abstract Factory
 - Delegate: This is a software design pattern where object composition is used to simulate multiple inheritances in order to take advantage of methods offered by different classes without inheriting them. The structure of the delegate pattern is as follows, Class A is composed of an object from Class B and provides an interface such that the methods from Class B can be called upon by functions in Class A.
 - Abstract Factory: This is a software design pattern that allows object creation through factory classes, without specifying what object is created and how it is created, deferring the creation behavior to be implemented inside the factory class that inherits from the abstract factory. This allows the creation of objects to be decoupled from the main program and separates the details of implementation of a set of objects from their general usage.
- Explain how you would approach implementing these two patterns in Ruby
 - Delegate: Used when a class uses another classes' function and uses the same name for themselves. Class A will be created in Ruby as normal, with given parameters and a method name (e.g. `method1(self, parameter)`). Class B will have a method also called “`method1(class A object)`” and return the result of calling the same method on that class A object
 - Abstract Factory: Used when we need to create multiple factory classes but also ensure that they are all using the same functions. We will create an abstract factory class that contains all general matrix functionality. Then we will create several matrix factories that inherits from our abstract factory but implements each method differently. Finally, in our sparse matrix module, we can take advantage of polymorphism to write general function calls for each matrix type.
- Are these patterns applicable to this problem? Explain your answer! (HINT: The answer is yes)
 - The reason we require the *Delegate Design Pattern*:
 - This pattern will be helpful for performing operations on a matrix or multiple matrices from a common class each with their own

implementation that specializes for its storage type. This lets us write a generic matrix module by taking advantage of polymorphism and composition.

■ The reason we require the *Abstract Factory Pattern*:

- Similarly, this pattern will aid us in generalizing our matrix module as we will only need to call on functions inside the abstract factory and then rely on polymorphism to execute the method relevant to that matrix type.

10. What implementation approach are you using (reuse class, modify class, inherit from class, compose with class, build new standalone class); justify your selection.

- Reuse class: This implementation approach will be primarily used as Ruby has several linear algebra libraries. The two most important are the standard Ruby Matrix library and NMatrix. We can reuse the classes in this module to reduce the amount of code we need to write.
- Modify class: This will be used when implementing functionality for tridiagonal matrices as most linear algebra libraries do not support this special case. This can also be used where existing libraries do not provide optimization for sparse matrices.
- Compose with class: This will be used for the delegate design pattern to call on different matrix operations of different types
- Inherit from class: This will be used with the abstract factory design pattern to construct our matrix factories.

11. Are iterators a good technique for sparse matrix manipulation? Are “custom” iterators required for this problem? (HINT: The answer is yes)

- Yes, we will need custom iterators as sparse matrices are usually not stored in conventional matrix format such as a NxN array. A custom iterator would allow us to access specific elements such as, elements on the main diagonal, elements in a specific column or elements in a specific row while hiding the details of the matrix storage implementation from the user.

12. What exceptions can occur during the processing of sparse matrices? And how should the system handle them?

- Sparse Matrix is no longer sparse: This depends on the context of the situation. For example, an addition operation would not care if the output is no longer sparse, but if a non-sparse matrix is given as input this would be an issue. The solution to this will be to write a method that determines the sparsity of a matrix.
- Matrix operations with incorrect dimensions sizes: Assertions should test dimensions before performing matrix operations and raise an exception if they do not match.
- Matrix too large exception: Assertions should test dimensions before performing matrix operations and raise an exception if they are too large.

- Inverting a matrix when determinant = 0: Assertions should test determinant before performing inverse operation and raise an exception if determinant = 0.
- Matrix shape exceptions, which could be thrown when an operation is called on a non-square matrix, that requires the matrix to be square (Such as calculating the determinant)

13. What information does the system require to create a sparse matrix object? Remember you are building for a set of unknown customers – what will they want?

- The dimensions of the matrix.
- Different storage types of the matrix. eg) Yale, List of lists or Dense
- Access to all elements in the matrix.
- Factorizations of sparse matrices
- Element-wise Operations on sparse matrices
- Matrix operations
- Easy to use API
- Creation methods
- Support for many data types
- Note: We should also be able to initialize an empty matrix.

14. What are the important quality characteristics of a sparse matrix package? Reusability? Efficiency? Efficiency of what?

- The efficiency of storage and data access, which we can achieve by ignoring zero elements.
- Usability of the matrix library. There are two quality of life characteristics that will enhance the usability of our library and make code written with our library easier to read:
 - Standard naming convention of each function.
 - Verbose exception handling for better reliability
- Reusability, Extensibility, Maintainability
 - When other developers implement the library, the use of consistent naming conventions will make our matrix library easier to use/work with.
- The efficiency of operations. If we do not consider the zero values of sparse matrices we can optimize the speed of our calculations.